

# Programming Assignment 4: Software Defined Networking

## Overview

For this project you will implement two control application for a software defined network (SDN). A layer-3 routing application will install rules in SDN switches to forward traffic to hosts using the shortest, valid path through the network. A distributed load balancer application will redirect new TCP connections to hosts in a round-robin order.

[Step 1: Getting Started](#)

[Step 2: Layer-3 Routing Application](#)

[Step 3: Distributed Load Balancer Application](#)

## Learning Objectives

After completing this programming assignment, students should be able to:

- Contrast SDN applications and traditional network control planes
- Create SDN applications that use proactive or reactive flow installation

## Clarifications

- You should make the following edit to the L3Routing.java file we provided:  
1) Change the line `this.knownHosts.remove(host)` in the `deviceRemoved(...)` event handler to `this.knownHosts.remove(device)`

## Step 1: Getting Started

Your SDN applications will be written in Java and run atop the [Floodlight](#) OpenFlow controller. You will use [Mininet](#) to emulate a variety of network topologies consisting of OpenFlow switches and hosts. You should continue to use your Mininet virtual machine (VM) for this project.

## Preparing Your Environment

Before beginning this project, there are some additional steps you need to complete to prepare your VM:

1. Install required packages

```
sudo apt-get update
sudo apt-get install -y curl traceroute ant openjdk-7-jdk git iputils-arping
```

2. Download Floodlight:

```
cd ~
git clone https://bitbucket.org/sdnhub/floodlight-plus.git
```

3. Get the starter codes and

```
cd ~/progAssign4
```

4. Symlink Floodlight

```
cd ~/progAssign4/
ln -s ../floodlight-plus
```
5. Patch Floodlight

```
cd ~/progAssign4/floodlight-plus
patch -p1 < ~/progAssign4/floodlight.patch
```
6. Install new version of Mininet and patch it

```
cd ~/progAssign4
./fixmininet.sh
```

If you want to use Eclipse (on a machine other than your VM) for development, you should do the following:

1. `cd floodlight-plus`  
`ant eclipse`
2. Open Eclipse and create a new workspace. **Do not** choose the floodlight-plus or progAssign4 folders as the location for your workspace.
3. Select *File* → *Import* → *General* → *Existing Projects into Workspace*. Click *Next*.
4. From *Select root directory* click *Browse*. Select the parent folder containing your floodlight-plus and progAssign4 folders.
5. Check the boxes for floodlight-plus and progAssign4. Click *Finish*.
6. Exclude the src/test/java folder from the build path:
  1. Right click on this folder in the *Project Explorer* pane.
  2. Select *Build Path* → *Remove from Build Path*.
7. Add the floodlight-plus project to the build path of progAssign4:
  1. Right click progAssign4 in the *Project Explorer* pane.
  2. Select *Build Path* → *Configure Build Path*.
  3. Under the *Projects* tab; click *Add*.
  4. Check the floodlight-plus project; click OK.

## Running Your Control Applications

1. Compile Floodlight and your applications:

```
cd ~/progAssign4/
ant
```

This will produce a jar file FloodlightWithApps.jar that includes the compiled code for Floodlight and your SDN applications.

2. Start Floodlight and your SDN applications:

```
java -jar FloodlightWithApps.jar -cf l3routing.prop
```

The above command will start Floodlight and only your layer-3 routing application. You can start both your layer-3 routing and load balancer applications by using loadbalancer.prop for the -cf (configuration file) argument.

When Floodlight starts, you should see output like the following:

```
19:15:37.239 INFO [n.f.c.m.FloodlightModuleLoader:main] Loading modules from file
l3routing.prop
19:15:37.666 INFO [n.f.c.i.Controller:main] Controller role set to MASTER
```

```

19:15:37.676 INFO [n.f.c.i.Controller:main] Flush switches on reconnect -- Disabled
19:15:37.697 INFO [ArpServer:main] Initializing ARPServer...
19:15:37.697 INFO [L3Routing:main] Initializing L3Routing...
19:15:38.835 INFO [n.f.l.i.LinkDiscoveryManager:main] Setting autoportfast feature to
OFF
19:15:38.891 INFO [ArpServer:main] Starting ARPServer...
19:15:38.892 INFO [L3Routing:main] Starting L3Routing...
19:15:39.122 INFO [o.s.s.i.c.FallbackCCProvider:main] Cluster not yet configured; using
fallback local configuration
19:15:39.122 INFO [o.s.s.i.SyncManager:main] [32767] Updating sync configuration
ClusterConfig [allNodes={32767=Node [hostname=localhost, port=6642, nodeId=32767,
domainId=32767]}, authScheme=NO_AUTH, keyStorePath=null, keyStorePassword is
unset]
19:15:39.185 INFO [o.s.s.i.r.RPCService:main] Listening for internal floodlight RPC on
localhost/127.0.0.1:6642
19:15:39.421 INFO [n.f.c.i.Controller:main] Listening for switch connections on
0.0.0.0/0.0.0.0:6633

```


Keep the terminal with Floodlight open, as you will need to see the output for debugging. Use another terminal for the next step. Now is a time to use [screen](#). You can install screen in your VM by running: `sudo apt-get install screen`.

### 3. Start Mininet:

```
sudo python ./run_mininet.py mesh,5
```

The above command will create a topology with 5 SDN switches and 5 hosts:

mesh,n: a complete graph with n switches and one host attached to each switch; for example, mesh,5 produces the following topology:



Once mininet has started, you should see Floodlight produce output like the following:

```

11:04:44.960 INFO [n.f.c.i.OFChannelHandler:New I/O server worker #2-1] New switch
connection from /127.0.0.1:50921
11:04:44.999 INFO [n.f.c.i.OFChannelHandler:New I/O server worker #2-1] Disconnected
switch [/127.0.0.1:50921 DPID[?]]
11:04:46.278 INFO [n.f.c.i.OFChannelHandler:New I/O server worker #2-2] New switch
connection from /127.0.0.1:50922
11:04:46.355 INFO [n.f.c.i.OFChannelHandler:New I/O server worker #2-2] Switch
OFSwitchBase [/127.0.0.1:50922 DPID[00:00:00:00:00:00:00:01]] bound to class class
net.floodlightcontroller.core.internal.OFSwitchImpl, writeThrottle=false, description
OFDescriptionStatistics [Vendor: Nicira, Inc., Model: Open vSwitch, Make: None,
Version: 2.0.1, S/N: None]
11:04:46.364 INFO [n.f.c.OFSwitchBase:New I/O server worker #2-2] Clearing all flows
on switch OFSwitchBase [/127.0.0.1:50922 DPID[00:00:00:00:00:00:00:01]]
11:04:46.377 WARN [n.f.c.i.C.s.notification:main] Switch 00:00:00:00:00:00:00:01
connected.

```

```
11:04:46.383 INFO [L3Routing:main] Switch s1 added
11:04:46.432 INFO [L3Routing:Topology Updates] Link s1:0 -> host updated
11:04:46.484 INFO [L3Routing:Topology Updates] Link s1:2 -> host updated
11:04:46.489 WARN [n.f.c.i.C.s.notification:main] Switch 00:00:00:00:00:00:01 port
s1-eth1 changed: UP
11:04:46.585 INFO [L3Routing:Topology Updates] Link s1:1 -> host updated
```

4. You can now run commands (e.g., ping) in mininet like you did for the previous assignments. However, until you complete Step 2, no packets will be forwarded.

You should **always** start Floodlight and your SDN applications **before** starting mininet. Also, we **recommend** that you restart Floodlight and your SDN applications whenever you restart mininet.

## Resources

If you have questions about SDN, OpenFlow, or Floodlight you may want to consult:

- [OpenFlow 1.3 Standard](#) — sections 2, 3, and 5.1 - 5.4 are likely to be the most useful
- [Floodlight-plus Javadoc](#)

## Step 2: Layer-3 Routing Application

### Background

In Programming Assignment 3, you wrote a router control plane that used RIP to construct a route table. The route table was then used by your router's data plane to determine where (i.e., out which interface) a packet should be forwarded.

Your first SDN application will also construct route tables. However, rather than exchanging RIP messages between routers, your application will construct route tables based on a global view of the network topology. The appropriate route table will be installed in each SDN switch by your application, and each SDN switch will forward packets according to the installed route table.

It is important to note that SDN switches do not directly support the route table structure used in the last programming assignment. Rather, SDN switches use a more general *flow table* structure, which can accommodate route tables (used in traditional layer-3 routers), MAC learning tables (used in traditional layer-2 switches), and many other constructs. Each entry, or *rule*, in a flow table has *match criteria* that defines (on the basis of fields in Ethernet, IP, TCP, UDP, and other headers) which packets the rule applies to. Each entry also has one or more *instructions/actions* which should be taken for each packet that matches the rule. Your layer-3 routing application will install entries that match packets based on their destination IP address (and Ethernet type), and execute an *output* action to send the packet out a specific port on the SDN switch.

(You'll use other match criteria and additional instructions/actions for the other SDN application you write, which is described in [Step 3](#).) The match criteria serves the same purpose as the *destination* and *mask* fields in a traditional route table, and the output action serves the same purpose as the *interface* field in a traditional route table. There is no concept of a *gateway* in SDN flow tables, but that's okay—your router only used the gateway to determine how to rewrite a packet's destination MAC address to ensure correct layer-2 forwarding, and we aren't using traditional layer-2 forwarding in SDN.

## Code Overview

The code for your layer-3 routing application will reside in `L3Routing.java` in the `cs.sdn.apps.l3routing` package. The file we provided already contains code to:

- Access host and topology information from other modules (or applications) included with Floodlight—see the `getHosts()`, `getSwitches()`, and `getLinks()` methods
- Receive notifications about changes in the network—`deviceAdded(...)`, `deviceRemoved(...)`, `deviceMoved(...)`, `switchAdded(...)`, `switchRemoved(...)`, and `linkDiscoveryUpdate(...)`

We have also provided code in the `cs.sdn.apps.utils` package for:

- An SDN application that responds to ARP requests from hosts—see `ArpServer.java`
- Telling a switch to install a rule in the flow table, remove rules from the flow table, and send a packet—see `SwitchCommands.java`

## TODOs

You need to complete the TODOs in `L3Routing.java` to install and remove flow table entries from SDN switches such that traffic is forwarded to a host using the shortest path. Note that Floodlight creates only one instance (i.e., object) of the `L3Routing` class; this instance handles all switches.

## Computing Shortest Paths

You should use the Bellman-Ford algorithm to compute the shortest paths to reach a host  $h$  from every other host  $h' \in H$ ,  $h \neq h'$  ( $H$  is the set of all hosts). You can use the `getHosts()`, `getSwitches()`, and `getLinks()` methods to get the topology information that you need to provide as input to the Bellman-Ford algorithm.

There will be two link objects between pairs of switches, one in each direction. Due to the way links are discovered, there may be a short period of time (tens of milliseconds) where the controller has a link object only in one direction. If you have a link object in one direction, you can assume the physical link is bidirectional.

When a host joins the network, both the `deviceAdded(...)` and `linkDiscoveryUpdate(...)` event handlers will be called. There are no guarantees on which order these event handlers are called. Thus, a host may be added but we may not yet know which switch it is linked to. The `isAttachedToSwitch()` method in the `Host` class will return true if we know the switch to which a host is connected, otherwise it will return false. If the method returns false, then you do not need to install any rules to route traffic to this host.

You can assume the following will always hold true in the network:

- The network is a connected graph. In other words, there will always be at least one possible path between every pair of switches.
- There is only one physical link between a pair of switches.
- Links are undirected. However, be aware that Floodlight maintains a `Link` object for each direction (i.e., there are two `Link` objects for each physical link).

## Installing Rules

Once you have determined the shortest path to reach host  $h$  from  $h'$ , you must install a rule in the flow table in every switch in the path. The rule should match IP packets (i.e., Ethernet

type is IPv4) whose destination IP is the IP address assigned to host *h*. You can specify this in Floodlight by creating a new OFMatch object and calling the set methods for the appropriate fields; you **must** set the Ethernet Type before you set the destination IP. The rule's action should be to output packets on the appropriate port in order to reach the next switch in the path. You can specify this in Floodlight by creating an OFInstructionApplyActions object whose set of actions consists of a single OFActionOutput object with the appropriate port number.

SDN switches have multiple flow tables—we discuss this more in [Step 3](#). For now, you should install rules in the table specified in the table class variable in the L3Routing class. Also, your rules should never timeout and have a default priority (both defined as constants in the SwitchCommands class).

## Testing and Debugging (You must show this to TA face-to-face on time!)

You should test your code by sending traffic between various hosts in the network topology—mininet's built-in pingall command is very useful for this.

To help you debug, you can view the contents of an SDN switch's flow tables by running the following command in your mininet VM (**not** in mininet itself):

```
sudo ovs-ofctl -O OpenFlow13 dump-flows s1
```

This will output the contents of s1's flow tables. Change the last argument to output the flow tables from a different switch.

## Triggering Event Handlers

- You can trigger the linkDiscoveryUpdate(...) event handler by running any of the following commands **in mininet** (substituting switch and host names as desired):
  - link s1 s2 down — takes down the link between s1 and s2; you can assume the network is a connected graph, so you should never take down a link that would result in a disconnected graph
  - link s1 s2 up — brings up the link between s1 and s2
  - link s1 h1 down — takes down the link between s1 and h1; this will also result in a deviceRemoved(...) event and the isAttachedToSwitch() method for the Host object for h1 will now return false
  - link s1 h1 up — brings up the link between s1 and h1; this will also result in a deviceMoved(...) event and the isAttachedToSwitch() method for the Host object for h1 will now return true
- You can trigger the deviceRemoved(...) event handler by taking down a link between a switch and a host, as described above
- You can trigger the deviceMoved(...) event handler by bringing up a link between a switch and a host, as described above
- You can trigger the switchRemoved(...) event handler by running the following command in a regular terminal window (**not** in mininet):

```
sudo ovs-vsctl del-br s1
```

Note that once a switch is removed, you cannot easily add it back without restarting mininet. You can assume the network is a connected graph, so you should never remove a switch that would result in a disconnected graph.

When you finish this, your screen may be like this:

```
0.6->00:00:00:00:00:00
14:57:16.107 INFO [ArpServer:New I/O server worker #2-1] Received ARP request fo
r 10.0.0.5 from 00:00:00:00:00:00
14:57:16.109 INFO [ArpServer:New I/O server worker #2-1] Sending ARP reply 10.0.
0.5->00:00:00:00:00:00
14:57:16.109 INFO [SwitchCommands:New I/O server worker #2-1] Forwarding packet:
arp
dl_vlan: untagged
dl_vlan_pcp: 0
dl_src: 00:00:00:00:00:05
dl_dst: 00:00:00:00:00:06
nw_src: 10.0.0.5
nw_dst: 10.0.0.6
14:57:16.110 INFO [SwitchCommands:New I/O server worker #2-2] Forwarding packet:
arp
dl_vlan: untagged
dl_vlan_pcp: 0
dl_src: 00:00:00:00:00:06
dl_dst: 00:00:00:00:00:05
nw_src: 10.0.0.6
nw_dst: 10.0.0.5
xx@xx-VirtualBox:~$ sudo ovs-ofctl -O OpenFlow13 dump-flows s1
OFPST_FLOW reply (OF1.3) (xid=0x2):
cookie=0x0, duration=25.322s, table=0, n_packets=0, n_bytes=0, priority=1,ip,nw
_dst=10.0.0.4 actions=output:4
cookie=0x0, duration=21.794s, table=0, n_packets=0, n_bytes=0, priority=1,ip,nw
_dst=10.0.0.6 actions=output:6
cookie=0x0, duration=21.659s, table=0, n_packets=0, n_bytes=0, priority=1,ip,nw
_dst=10.0.0.1 actions=output:1
cookie=0x0, duration=21.705s, table=0, n_packets=0, n_bytes=0, priority=1,ip,nw
_dst=10.0.0.5 actions=output:5
cookie=0x0, duration=21.750s, table=0, n_packets=0, n_bytes=0, priority=1,ip,nw
_dst=10.0.0.3 actions=output:3
xx@xx-VirtualBox:~$
h3 -> h1 h2 h4 h5 h6
h4 -> h1 h2 h3 h5 h6
h5 -> h1 h2 h3 h4 h6
h6 -> h1 h2 h3 h4 h5
*** Results: 0% dropped (30/30 received)
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3 h4 h5 h6
h2 -> h1 h3 h4 h5 h6
h3 -> h1 h2 h4 h5 h6
h4 -> h1 h2 h3 h5 h6
h5 -> h1 h2 h3 h4 h6
h6 -> h1 h2 h3 h4 h5
*** Results: 0% dropped (30/30 received)
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3 h4 h5 h6
h2 -> h1 h3 h4 h5 h6
h3 -> h1 h2 h4 h5 h6
h4 -> h1 h2 h3 h5 h6
h5 -> h1 h2 h3 h4 h6
h6 -> h1 h2 h3 h4 h5
*** Results: 0% dropped (30/30 received)
mininet>
```

## Step 3: Distributed Load Balancer Application

### Background

Networks employ *load balancing* to distribute client requests among a collection of hosts running a specific service (e.g., a web server). In class, we discussed how DNS could be used to implement load balancing. Load balancing is also commonly implemented using a special piece of hardware.

A hardware load balancer is placed in the network and configured with an IP address (e.g., 10.0.100.1) and a set of hosts among which it should distribute requests (e.g., 10.0.0.2 and 10.0.0.3). Clients wanting to communicate with a service (e.g., a web server) running on those hosts are provided with the IP address of the load balancer, not the IP address of a specific host. Clients initiate a TCP connection to the IP address of the load balancer (10.0.100.1) and the TCP port associated with the service (e.g., port 80).

For each new TCP connection, the load balancer selects one of the specified hosts (usually in round robin order). The load balancer maintains a mapping of active connections—identified by the client's IP and TCP port—to the assigned hosts.

For all packets sent from clients to the load balancer, the load balancer rewrites the destination IP and MAC addresses to the IP and MAC addresses of the selected host. The mapping information stored by the load balancer is used to determine the appropriate host IP and MAC addresses that should be written into a packet arriving from a client. For all packets sent from servers to clients, the load balancer rewrites the source IP and MAC addresses to the IP and MAC addresses of the load balancer.

Your second SDN application will implement the same functionality as a set of hardware load balancers. Your application will be provided with a list of *virtual IPs* and a set of hosts among which connections to the virtual IPs should be load balanced. (We use the term *virtual IP* because the IP address is not actually assigned to any node in the network.) When clients initiate TCP connections with a specific *virtual IP*, SDN switches will send the TCP SYN packet to the SDN controller. Your SDN application will select a host from a pre-defined set, and install rules in an SDN switch to rewrite the IP and MAC addresses of packets associated with the connection. You will also instruct the SDN switch to match the modified packets



against the flow rules installed by your layer-3 routing application and apply the appropriate actions (i.e., send the packets out the appropriate ports).

## Code Overview

The code for your load balancer application will reside in `LoadBalancer.java` in the `cs.sdn.apps.loadbalancer` package. The file we provided already contains code to:

- Receive a notification when a switch joins the network—`switchAdded(...)`
- Receive a packet from a switch when the packet did not match any entries in the switch's flow table—`receive(...)`

The `LoadBalancerInstance` class represents a single distributed load balancer. (We use the term *distributed* because the load balancing is performed at many switches, rather than at a single hardware load balancer.) Each load balancer instance has a virtual IP address, virtual MAC address, and set of hosts among which TCP connections should be distributed. The instances class variable in the `LoadBalancer` class maps a virtual IP address to a specific load balancer instance.

## TODOs

You need to complete the TODOs in `LoadBalancer.java` to:

- Install rules in every switch (when it joins the network) to:
  - Notify the controller when a client initiates a TCP connection with a virtual IP—we cannot specify TCP flags in match criteria, so the SDN switch will notify the controller of each TCP packet sent to a virtual IP which did not match a connection-specific rule (described below)
  - Notify the controller when a client issues an ARP request for the MAC address associated with a virtual IP
  - Match all other packets against the rules in the next table in the switch (described [below](#))
- Install connection-specific rules for each new connection to a virtual IP to:
  - Rewrite the destination IP and MAC address of TCP packets sent from a client to the virtual IP
  - Rewrite the source IP and MAC address of TCP packets sent from server to client

Connection-specific rules should match packets on the basis of Ethernet type, source IP address, destination IP address, protocol, TCP source port, and TCP destination port. Connection-specific rules should take precedence over the rules that send TCP packets to the controller, otherwise every TCP packet would be sent to the controller. Therefore, these rules should have a higher priority than the rules installed when a switch joins the network. Also, we want connection-specific rules to be removed when a TCP connection ends, so connection-specific rules should have an idle timeout of 20 seconds.

- Construct and send an ARP reply packet when a client requests the MAC address associated with a virtual IP



## Multiple Tables

Your load balancer application should work in tandem with your layer-3 routing application. To achieve this, you will need to leverage the *multiple tables* feature of OpenFlow switches. When packets first arrive at an OpenFlow switch, they are matched against the rules in table 0. The actions for these rules can specify that the packets be modified, output, sent to the controller, and/or matched against the rules in a different table.

Your load balancer application should install rules in the table specified in the table class variable in the LoadBalancer class—set to table 0 in the loadbalancer.prop configuration file. The connection-specific rules that modify IP and MAC addresses should include an [instruction](#) to match the modified packets against the rules installed by your layer-3 routing application. Since your layer-3 routing application will install rules in the table class variable in the L3Routing class, this instruction should direct packets to the table defined in this class variable. The modified packet will then be matched against these rules and forwarded out the appropriate port.

All packets which are not TCP packets destined for a virtual IP, or packets associated with a connection that has already been assigned to a specific host, should be sent directly to the table used by your layer-3 routing application.

## Sending ARP Packets

When a client wants to initiate a connection with the virtual IP, it will need to determine the MAC address associated with the virtual IP using ARP. The client does not know the IP is virtual, and since it's not actually assigned to any host, your SDN application must take responsibility for replying to these requests.

You can construct an ARP reply packet using the classes in the net.floodlightcontroller.packet package. You can use the sendPacket(...) method in the SwitchCommands class to send the packet.

## Rule Instructions/Actions

When a rule should send a packet to the controller, the rule should include an OFInstructionApplyActions whose set of actions consists of a single OFActionOutput with OFPort.OFPP\_CONTROLLER as the port number.

When a rule should rewrite the destination IP and MAC addresses of a packet, the rule should include an OFInstructionApplyActions whose set of actions consists of:

- An OFActionSetField with a field type of OFOXMFieldType.ETH\_DST and the desired MAC address as the value
- An OFActionSetField with a field type of OFOXMFieldType.IPV4\_DST and the desired IP address as the value

The actions for rewriting the source IP and MAC addresses of a packet are similar.

When a packet should be processed by the SDN switch based on the rules installed by your layer-3 routing application, a rule should include an OFInstructionGotoTable whose table number is the value specified in the table class variable in the L3Routing class.

## Testing and Debugging(You must show this to TA face-to-face on time!)

You should test your code by issuing web requests (using curl) from a client host to the virtual IPs.

You can add or remove virtual IPs and hosts by modifying the loadbalancer.prop file.

To see which packets a host is sending/receiving, run:

tcpdump -v -n -i hN-eth0  
replacing N with the host's number.  
When you finish this, your screen may be like this:

The image shows two terminal windows. The left window displays the output of running a Mininet script to create a network topology with 6 hosts (h1-h6) and 6 switches (s1-s6). It shows the configuration of links, starting of services on each host, and several curl commands being executed from host h1 to 10.0.100.1, with corresponding network statistics tables. The right window shows a detailed view of the network topology, listing the connections between switches and their distances.

```

xx@xx-VirtualBox:~/assign4$ sudo python ./run_mininet.py mesh,6
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4 h5 h6
*** Adding switches:
s1 s2 s3 s4 s5 s6
*** Adding links:
(h1, s1) (h2, s2) (h3, s3) (h4, s4) (h5, s5) (h6, s6) (s1, s2) (s1, s3) (s1, s4)
(s1, s5) (s1, s6) (s2, s3) (s2, s4) (s2, s5) (s2, s6) (s3, s4) (s3, s5) (s3, s6)
(s4, s5) (s4, s6) (s5, s6)
*** Configuring hosts
h1 h2 h3 h4 h5 h6
*** Starting controller
*** Starting 6 switches
s1 s2 s3 s4 s5 s6
*** ARPing from host h1
*** Starting SimpleHTTPServer on host h1
*** ARPing from host h2
*** Starting SimpleHTTPServer on host h2
*** ARPing from host h3
*** Starting SimpleHTTPServer on host h3
*** ARPing from host h4
*** Starting SimpleHTTPServer on host h4
*** ARPing from host h5
*** Starting SimpleHTTPServer on host h5
*** ARPing from host h6
*** Starting SimpleHTTPServer on host h6
*** Starting CLI:
mininet> h1 curl 10.0.100.1
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 23 100 23 0 0 21 0 0:00:01 0:00:01 --:--:-- 21
WEB PAGE SERVED BY: h2
mininet> h1 curl 10.0.100.1
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 23 100 23 0 0 21 0 0:00:01 0:00:01 --:--:-- 21
WEB PAGE SERVED BY: h3
mininet> h1 curl 10.0.100.1
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 23 100 23 0 0 22 0 0:00:01 0:00:01 --:--:-- 22
WEB PAGE SERVED BY: h2
mininet> h1 curl 10.0.100.1
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 23 100 23 0 0 22 0 0:00:01 0:00:01 --:--:-- 23
WEB PAGE SERVED BY: h3
mininet>
  
```

```

The other_sw: 2
DistToSwitch: 1 | DistToOther: 1
The other_sw: 4
DistToSwitch: 1 | DistToOther: 0
The other_sw: 6
DistToSwitch: 1 | DistToOther: 1
Processing ID: 3
The other_sw: 5
DistToSwitch: 1 | DistToOther: 1
The other_sw: 1
DistToSwitch: 1 | DistToOther: 1
The other_sw: 2
DistToSwitch: 1 | DistToOther: 1
The other_sw: 4
DistToSwitch: 1 | DistToOther: 0
The other_sw: 6
DistToSwitch: 1 | DistToOther: 1
ID: 6 | V: 1
ID: 5 | V: 1
ID: 2 | V: 1
ID: 1 | V: 1
ID: 3 | V: 1
ID: 4 | V: 0
  
```

## Submission Instructions

You must submit a single tar file of the src directory containing the Java source files for your SDN applications. Please submit the entire src directory; do not submit any other files or directories. Also, you should submit a readme file.  
You must submit a single tar file named assign4\_username1\_username2 containing the above. Please submit only one tar file per group.

## Grading

1) Correctness 80%.

L3Routing 40

Distributed Load Balance 40

2) Code quality 20% (Dont copy from others! )

Make code more readable.

## Acknowledgements

We acknowledge CS 640: Computer Networks with University of Wisconsin-Madison.